

Treemaps, Heatmaps, Alluvials and Streamgraphs

- [Intro](#)
- [Heatmaps](#)
- [Treemaps](#)
- [An Exploration with a Heatmap of World Happiness](#)
- [Alluvials](#)
- [nycflights23 dataset to explore heatmaps](#)
- [Streamgraphs using babynames](#)
- [Homework Week 5](#)

Intro to Heatmaps

The following heatmap is one of my favorites created to show how climate change is affecting global temperatures. It tells a story of temperatures warming over a long period of time.

[Climate Change is Rewriting the History Books – There hasn't been a cool month in 628 months \(dated from 2015\)](#)

This week, you will learn how to make a heatmap in R from Flowing Data Tutorial

<https://flowingdata.com/2010/01/21/how-to-make-a-heatmap-a-quick-and-easy-solution/>



Heat Tree Stream Alluvial - Part 1

R Saidi

So many ways to visualize data

Load the packages and the data from flowingdata.com website

```
library(treemap)
library(tidyverse)
library(RColorBrewer)
library(webshot2)
```

Heatmaps

A heatmap is a literal way of visualizing a table of numbers, where you substitute the numbers with colored cells. There are two fundamentally different categories of heat maps: the cluster heat map and the spatial heat map. In a cluster heat map, magnitudes are laid out into a matrix of fixed cell size whose rows and columns are discrete categories, and the sorting of rows and columns is intentional. The size of the cell is arbitrary but large enough to be clearly visible. By contrast, the position of a magnitude in a spatial heat map is forced by the location of the magnitude in that space, and there is no notion of cells; the phenomenon is considered to vary continuously. (Wikipedia)

Load the nba data from Yau's website

This data appears to contain data about 2008 NBA player stats.

```
nba <- read.csv("http://datasets.flowingdata.com/ppg2008.csv")
#apparently you have to use read.csv here instead of read_csv
head(nba)
```

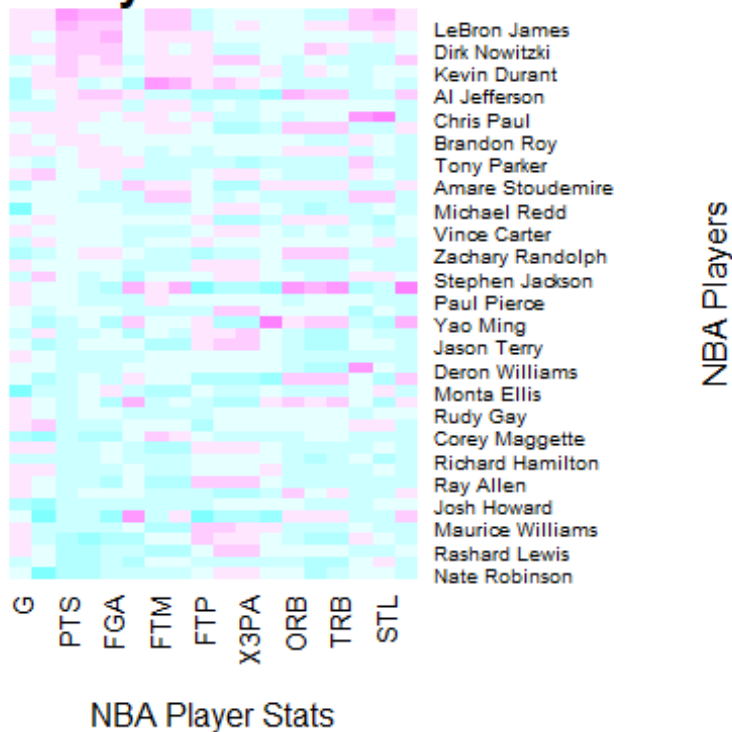
	Name	G	MIN	PTS	FGM	FGA	FGP	FTM	FTA	FTP	X3PM	X3PA	X3PP	ORB
1	Dwyane Wade	79	38.6	30.2	10.8	22.0	0.491	7.5	9.8	0.765	1.1	3.5	0.317	1.1
2	LeBron James	81	37.7	28.4	9.7	19.9	0.489	7.3	9.4	0.780	1.6	4.7	0.344	1.3
3	Kobe Bryant	82	36.2	26.8	9.8	20.9	0.467	5.9	6.9	0.856	1.4	4.1	0.351	1.1
4	Dirk Nowitzki	81	37.7	25.9	9.6	20.0	0.479	6.0	6.7	0.890	0.8	2.1	0.359	1.1
5	Danny Granger	67	36.2	25.8	8.5	19.1	0.447	6.0	6.9	0.878	2.7	6.7	0.404	0.7
6	Kevin Durant	74	39.0	25.3	8.9	18.8	0.476	6.1	7.1	0.863	1.3	3.1	0.422	1.0
	DRB	TRB	AST	STL	BLK	TO	PF							
1	3.9	5.0	7.5	2.2	1.3	3.4	2.3							
2	6.3	7.6	7.2	1.7	1.1	3.0	1.7							
3	4.1	5.2	4.9	1.5	0.5	2.6	2.3							
4	7.3	8.4	2.4	0.8	0.8	1.9	2.2							
5	4.4	5.1	2.7	1.0	1.4	2.5	3.1							
6	5.5	6.5	2.8	1.3	0.7	3.0	1.8							

Create a cool-color heatmap

This older heatmap function requires the data to be formatted as a matrix using the `data.matrix`

```
nba <- nba[order(nba$PTS),]
row.names(nba) <- nba$Name
nba <- nba[,2:19]
nba_matrix <- data.matrix(nba)
nba_heatmap <- heatmap(nba_matrix,
                        Rowv=NA,
                        Colv=NA,
                        col = cm.colors(10),
                        scale="column",
                        margins=c(5,10),
                        xlab = "NBA Player Stats",
                        ylab = "NBA Players",
                        main = "NBA Player Stats in 2008")
```

NBA Player Stats in 2008



What did that plot show?

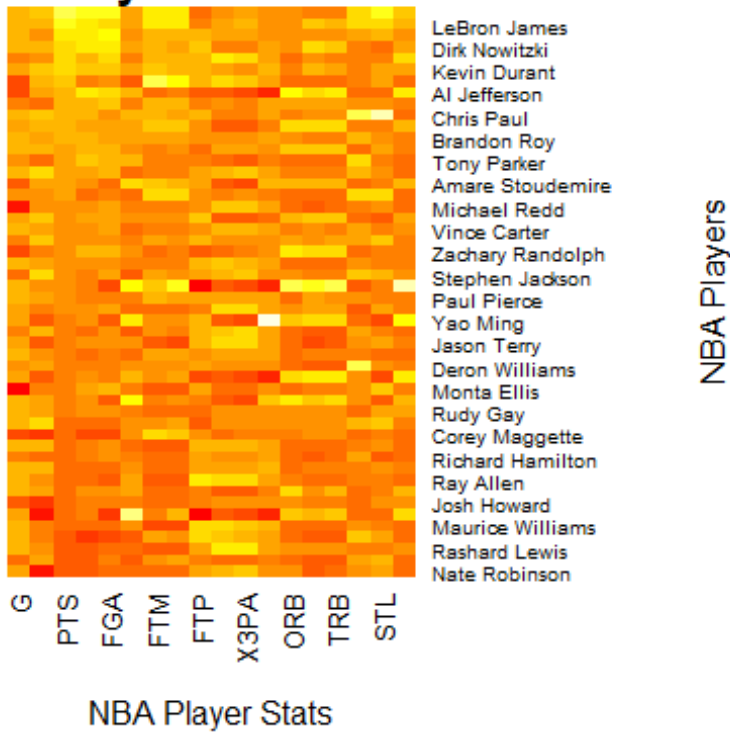
The basic layout of the heatmap relies on the parameters rows, columns and values. You can think of them like aesthetics in `ggplot2::ggplot()`, similar to something like `aes(x = columns, y = rows, fill = values)`.

Change to warm color palette

```
nba_heatmap <- heatmap(nba_matrix,
                        Rowv=NA,
                        Colv=NA,
                        col = heat.colors(20),
                        scale="column",
                        margins=c(5,10),
                        xlab = "NBA Player Stats",
```

```
ylab = "NBA Players",
main = "NBA Player Stats in 2008")
```

NBA Player Stats in 2008



Use the viridis color palette

For some reason the viridis colors from viridisLite package default to give dendrite clustering (the branches).

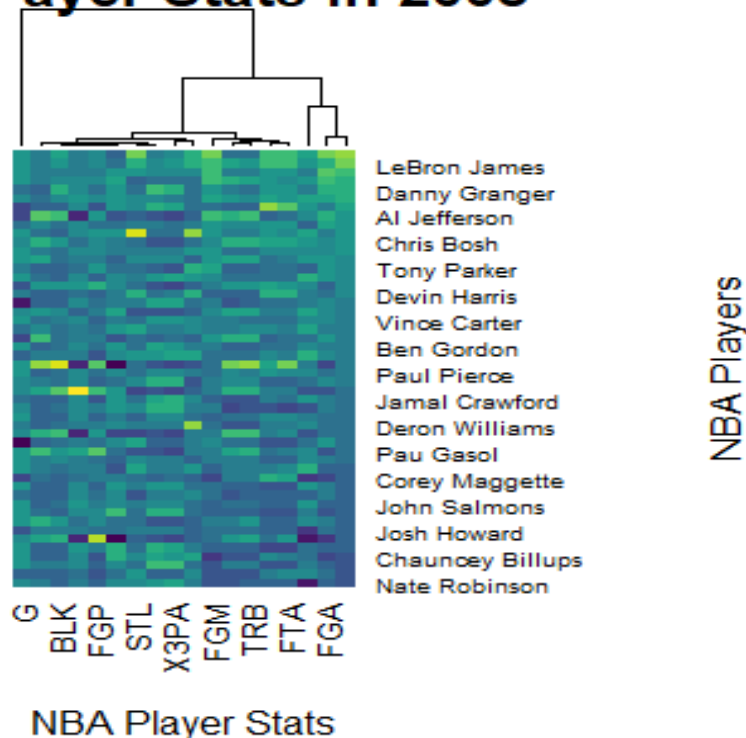
```
library(viridis)
```

Loading required package: viridisLite

```
# Loading required package: viridis
```

```
nba_heatmap <- heatmap(nba_matrix,
  Rowv=NA,
  col = viridis(20),
  scale="column",
  margins=c(5,10),
  xlab = "NBA Player Stats",
  ylab = "NBA Players",
  keep.dendro = FALSE,
  main = "NBA Payer Stats in 2008")
```

NBA Payer Stats in 2008



Treemaps

Treemaps display hierarchical (tree-structured) data as a set of nested rectangles. Each branch of the tree is given a rectangle, which is then tiled with smaller rectangles representing sub-branches. A leaf node's rectangle has an area proportional to a specified dimension of the data.[1] Often the leaf nodes are colored to show a separate dimension of the data.

When the color and size dimensions are correlated in some way with the tree structure, one can often easily see patterns that would be difficult to spot in other ways, such as whether a certain color is particularly relevant. A second advantage of treemaps is that, by construction, they make efficient use of space. As a result, they can legibly display thousands of items on the screen simultaneously.

The Downside to Treemaps

The downside of treemaps is that as the aspect ratio is optimized, the order of placement becomes less predictable. As the order becomes more stable, the aspect ratio is degraded. (Wikipedia)

Use Nathan Yau's dataset from the flowingdata website: <http://datasets.flowingdata.com/post-data.txt> You will need the package "treemap" and the package "RColorBrewer".

Create a treemap which explores categories of views

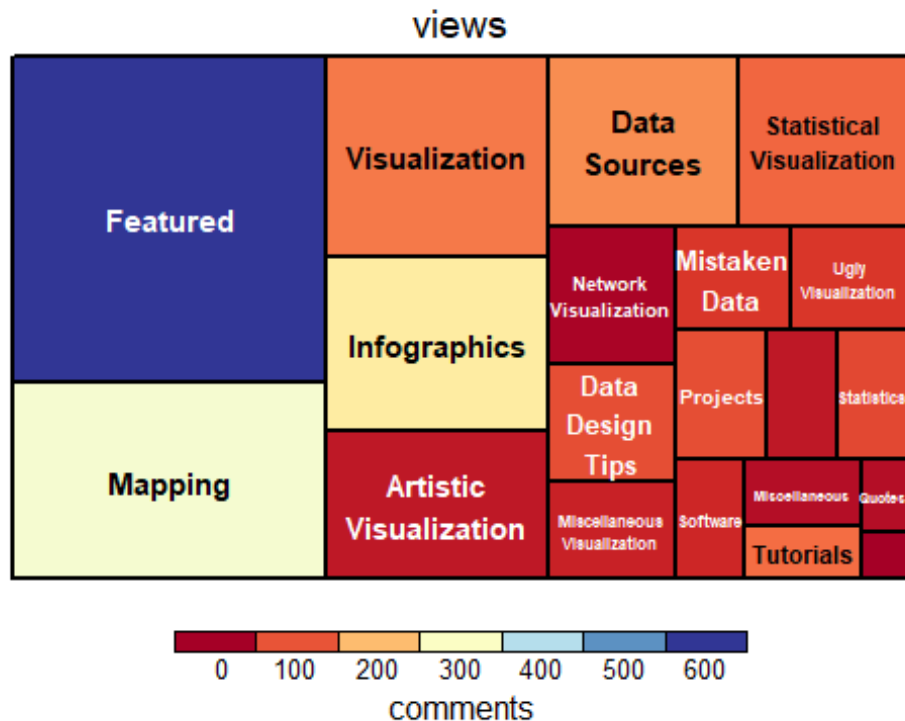
Load the data for creating a treemap from Nathan Yao's flowing data which explores number of views and comments for different categories of posts on his website.

```
flowingdata <- read.csv("http://datasets.flowingdata.com/post-data.txt")
# again, here use read.csv instead of read_csv
head(flowingdata)
```

	id	views	comments	category
1	5019	148896	28	Artistic Visualization
2	1416	81374	26	Visualization
3	1416	81374	26	Featured
4	3485	80819	37	Featured
5	3485	80819	37	Mapping
6	3485	80819	37	Data Sources

Use RColorBrewer to change the palette to RdYlBu

```
treemap(flowingdata, index="category", vSize="views",
        vColor="comments", type="manual",
        # note: type = "manual" changes to red yellow blue
        palette="RdYlBu")
```



Notice the following:

- The index is a categorical variable - in this case, "category" of post
- The size of the box is by number of views of the post
- The heatmap color is by number of comments for the post
- Notice how the treemap includes a legend for number of comments *

A heatmap of World Happiness

Set your working directory and read in the happiness19.csv from the class google drive.

```
setwd("C:/Users/rsaiddi/Dropbox/Rachel/MontColl/Datasets/Datasets")
happy19 <- read_csv("happiness2019.csv")
head(happy19)

# A tibble: 6 × 9
  `Overall rank` `Country or region` Score `GDP per capita` `Social support`
      <dbl> <chr>          <dbl>          <dbl>          <dbl>
1           1 Finland            7.77            1.34            1.59
2           2 Denmark            7.6             1.38            1.57
3           3 Norway             7.55            1.49            1.58
4           4 Iceland            7.49            1.38            1.62
5           5 Netherlands        7.49            1.40            1.52
6           6 Switzerland        7.48            1.45            1.53
# [i] 4 more variables: `Healthy life expectancy` <dbl>,
#   `Freedom to make life choices` <dbl>, Generosity <dbl>,
#   `Perceptions of corruption` <dbl>
```

We can see that there are 156 countries ranked by their “happiness score” based on other measurements.

Clean the happiness dataset to work with it

first remove the first column for “overall_rank”. Clean the remaining headers.

```
happy <- happy19 |>
  select(-`Overall rank`)
names(happy) <- tolower(names(happy))
names(happy) <- gsub(" ", "_", names(happy))
```

Because there are 156 countries, narrow the inclusion criteria to be for the top 20 scoring countries. Then do the same for the lowest 20.

```
happytop <- happy |>
  arrange(desc(score)) |>
  mutate(happy = "top") |> # add a column for use later
  head(20)
happytop

# A tibble: 20 × 9

happybottom <- happy |>
  arrange(score) |>
  mutate(happy = "bottom") |>
  head(20)
happybottom

# A tibble: 20 × 9
```

Then convert from wide to long format.

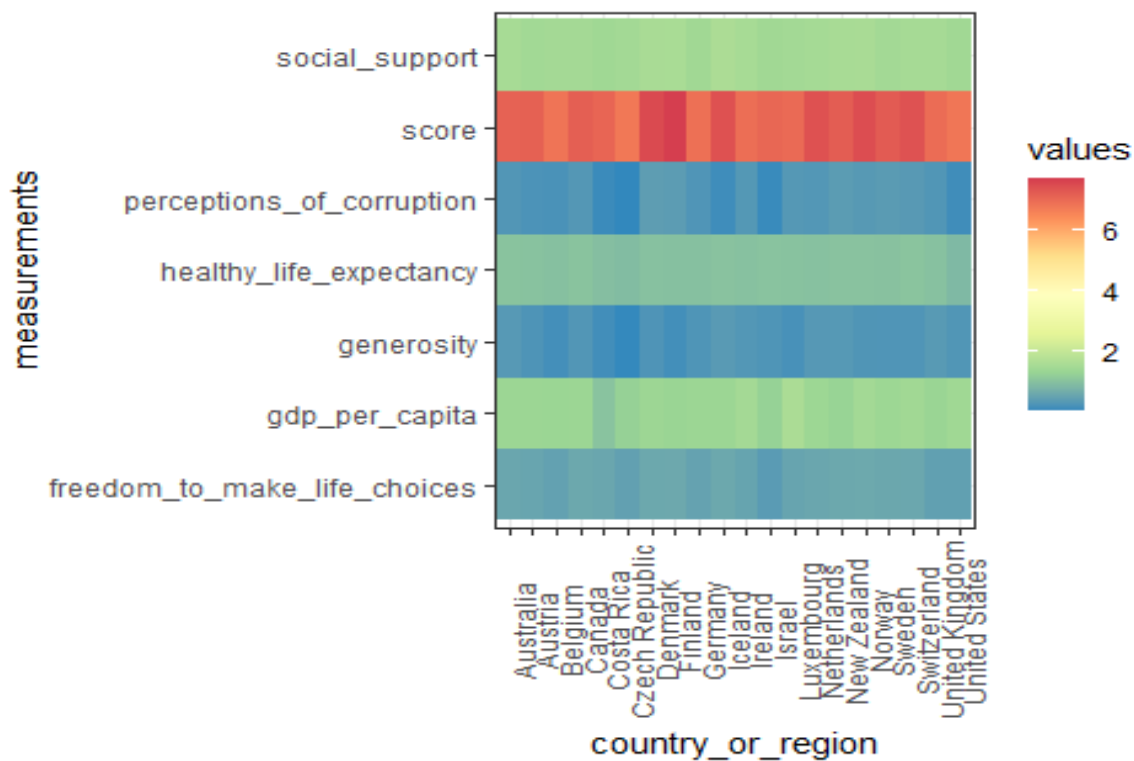
```
happy_longtop <- happytop |>
  pivot_longer(cols = 2:8,
               names_to = "measurements",
               values_to = "values")
happy_longtop

# A tibble: 140 × 4
  country_or_region happy measurements values
  <chr>             <chr> <chr>          <dbl>
1 Finland          top    score          7.77
2 Finland          top    gdp_per_capita 1.34
3 Finland          top    social_support 1.59
4 Finland          top    healthy_life_expectancy 0.986
5 Finland          top    freedom_to_make_life_choices 0.596
6 Finland          top    generosity      0.153
7 Finland          top    perceptions_of_corruption 0.393
8 Denmark          top    score          7.6
9 Denmark          top    gdp_per_capita 1.38
10 Denmark         top    social_support 1.57
# [i] 130 more rows

happy_longbottom <- happybottom |>
  pivot_longer(cols = 2:8,
               names_to = "measurements",
               values_to = "values")
happy_longbottom

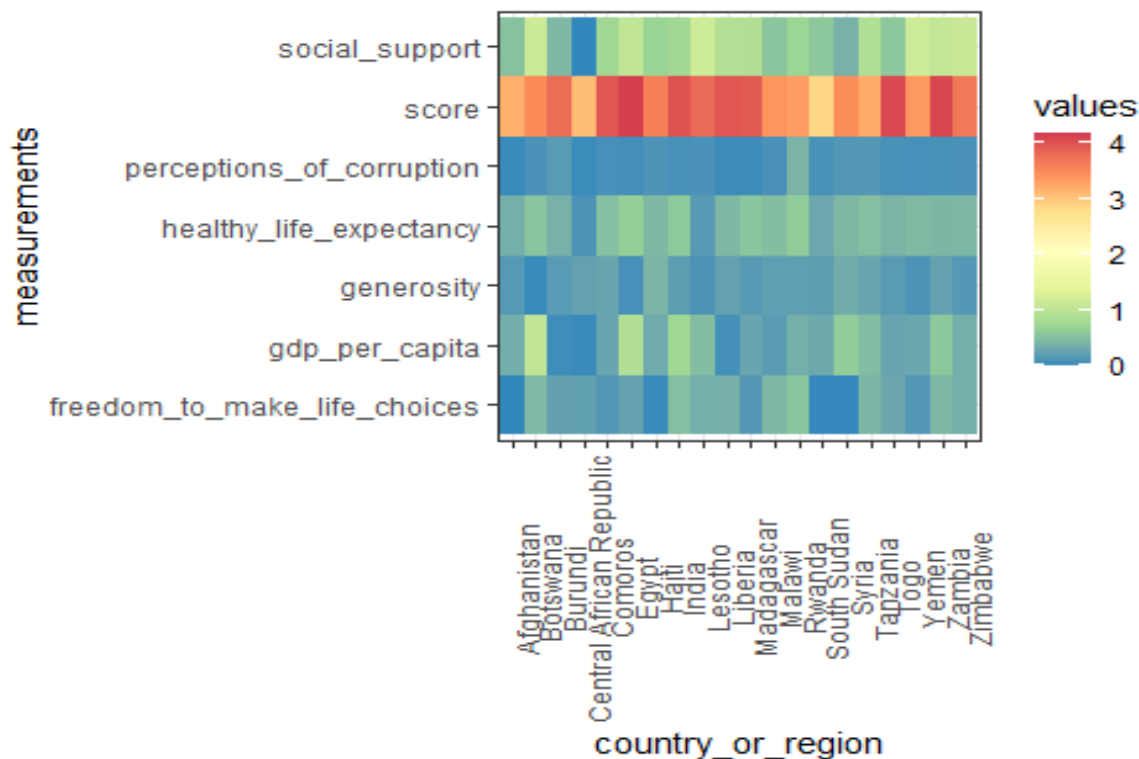
# A tibble: 140 × 4
  country_or_region happy measurements values
  <chr>             <chr> <chr>          <dbl>
1 South Sudan      bottom score          2.85
2 South Sudan      bottom gdp_per_capita 0.306
3 South Sudan      bottom social_support 0.575
4 South Sudan      bottom healthy_life_expectancy 0.295
5 South Sudan      bottom freedom_to_make_life_choices 0.01
6 South Sudan      bottom generosity      0.202
7 South Sudan      bottom perceptions_of_corruption 0.091
8 Central African Republic bottom score          3.08
9 Central African Republic bottom gdp_per_capita 0.026
10 Central African Republic bottom social_support 0
# [i] 130 more rows

ggplot(data = happy_longtop, aes(x=country_or_region, y=measurements, fill = values)) +
  geom_tile()+
  scale_fill_distiller(palette="Spectral") +
  theme_bw()+
  theme(axis.text.x = element_text(angle = 90))
```

What do we notice a common patterns and different patterns at the top?

```
ggplot(data = happy_longbottom, aes(x=country_or_region, y=measurements, fill = values))
+
  geom_tile()+
  scale_fill_distiller(palette="Spectral") +
  theme_bw()+
  theme(axis.text.x = element_text(angle = 90))
```



What do we notice a common patterns and different patterns at the bottom?

Put the top 20 and bottom 20 together to compare the two plots

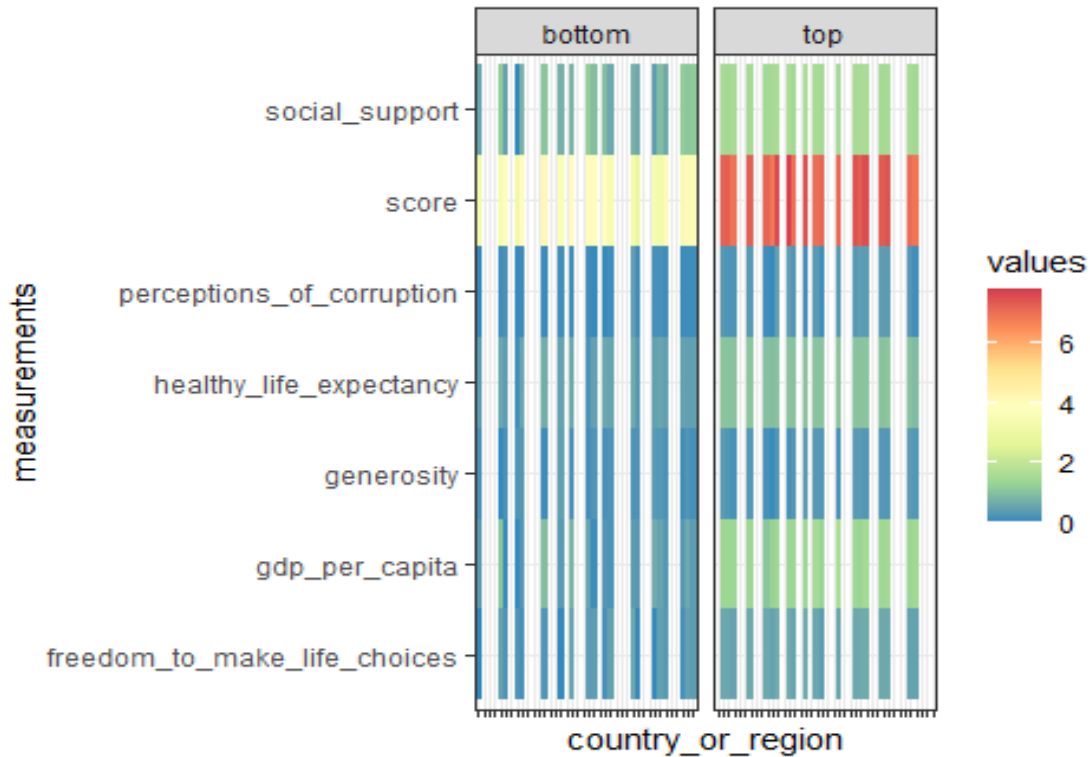
```
newdf <- rbind(happytop, happybottom)
newdf_long <- newdf |>
  pivot_longer(cols = 2:8,
               names_to = "measurements",
               values_to = "values")
head(newdf_long)
```

#	country_or_region	happy	measurements	values
	<chr>	<chr>	<chr>	<dbl>
1	Finland	top	score	7.77
2	Finland	top	gdp_per_capita	1.34
3	Finland	top	social_support	1.59
4	Finland	top	healthy_life_expectancy	0.986
5	Finland	top	freedom_to_make_life_choices	0.596
6	Finland	top	generosity	0.153

create a facet plot of the geom_tile

```
ggplot(data = newdf_long, aes(x=country_or_region, y=measurements, fill = values)) +
  geom_tile()+
  scale_fill_distiller(palette="Spectral") +
  facet_grid(~happy) +
  theme_bw()+
```

```
theme(axis.text.x = element_blank()) # remove the countries to generally compare top and bottom ranked countries
```



The facet plot places the top and bottom countries on the same scale, and now we can really make comparisons.

What do you notice now?

Alluvials

Load the alluvial package

Refugees is a prebuilt dataset in the alluvial package

If you want to save the prebuilt dataset to your folder, use the write_csv function

```
library(alluvial)
library(ggalluvial)
data(Refugees)
```

Show UNHCR-recognised refugees

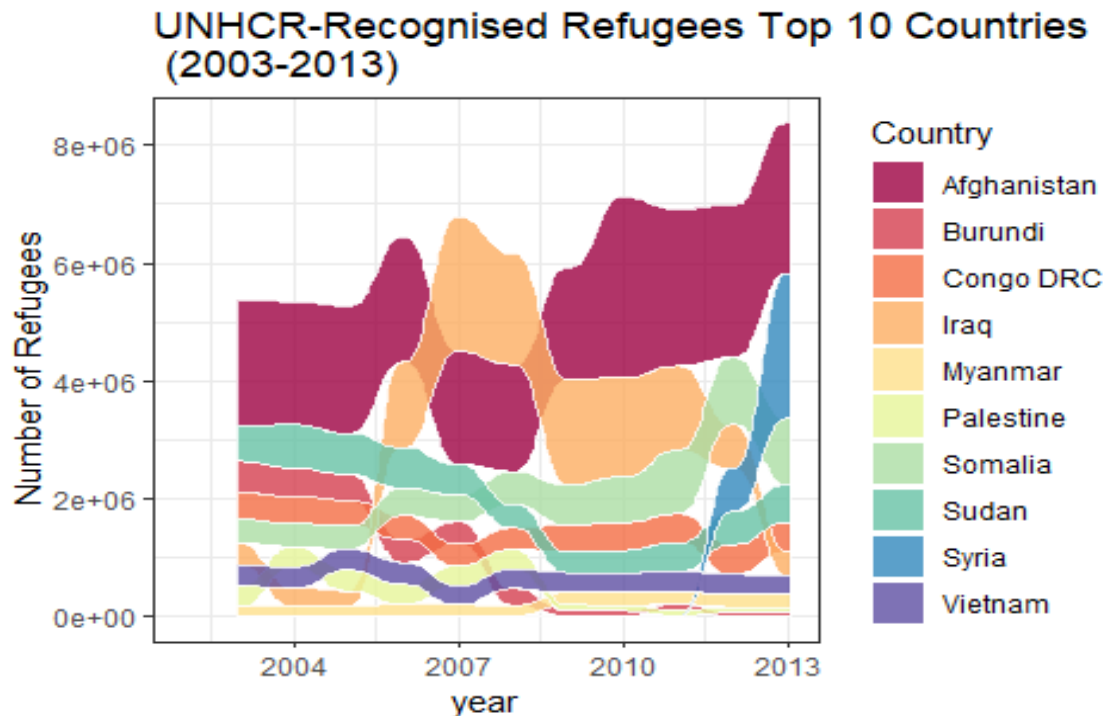
Top 10 most affected countries causing refugees from 2003-2013 Alluvials need the variables: *time-variable*, *value*, *category*

```
ggalluv <- Refugees |>
  ggplot(aes(x = year, y = refugees, alluvium = country)) +
  theme_bw() +
```

```
geom_alluvium(aes(fill = country),
              color = "white",
              width = .1,
              alpha = .8,
              decreasing = FALSE) +
scale_fill_brewer(palette = "Spectral") +
# Spectral has enough colors for all countries listed
scale_x_continuous(lim = c(2002, 2013)) +
labs(title = "UNHCR-Recognised Refugees Top 10 Countries\n (2003-2013)",
     # \n breaks the long title
     y = "Number of Refugees",
     fill = "Country",
     caption = "Source: United Nations High Commissioner for Refugees (UNHCR)")
```

Plot the Alluvial

ggalluv



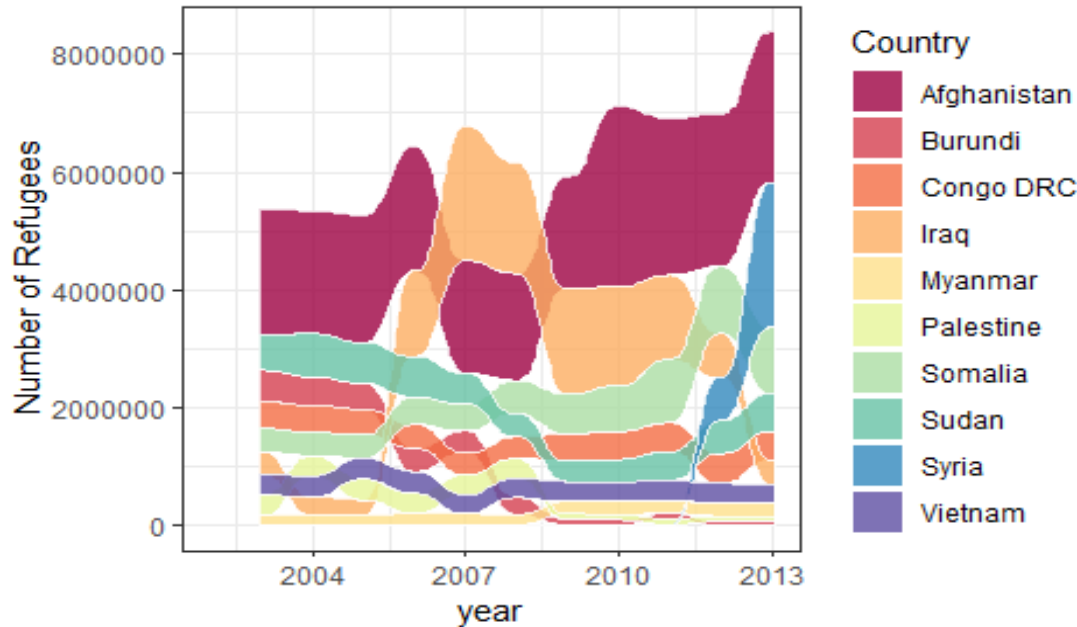
ce: United Nations High Commissioner for Refugees (UNHCR)

A final touch to fix the y-axis scale

Notice the y-values are in scientific notation. We can convert them to standard notation with options scipen function

```
options(scipen = 999)
ggalluv
```

UNHCR-Recognised Refugees Top 10 Countries (2003-2013)



Use the dataset NYCFlights23 to create a heatmap that explores Late Arrivals

Source: FAA Aircraft registry,
https://www.faa.gov/licenses_certificates/aircraft_certification/aircraft_registry/releasable_aircraft_download/

```
#install.packages("nycflights23")
library(nycflights23)
library(RColorBrewer)
data(flights)
data(airlines)
```

Create an initial scatterplot with loess smoother for distance to delays

Use "group_by" together with summarise functions

Remove observations with NA values from distand and arr_delay variables - notice number of rows changed from 336,776 to 327,346

Never use the function "na.omit"!!!!

```
flights_nona <- flights |>
  filter(!is.na(distance) & !is.na(arr_delay))
# remove na's for distance and arr_delay
```

Use group_by and summarise to create a summary table

The table includes, counts for each tail number, mean distance traveled, and mean arrival delay

```
by_tailnum <- flights_nona |>
  group_by(tailnum) |> # group all tailnumbers together
  summarise(count = n(), # counts totals for each tailnumber
            dist = mean(distance), # calculates the mean distance traveled
            delay = mean(arr_delay)
            ) # calculates the mean arrival delay
head(by_tailnum)

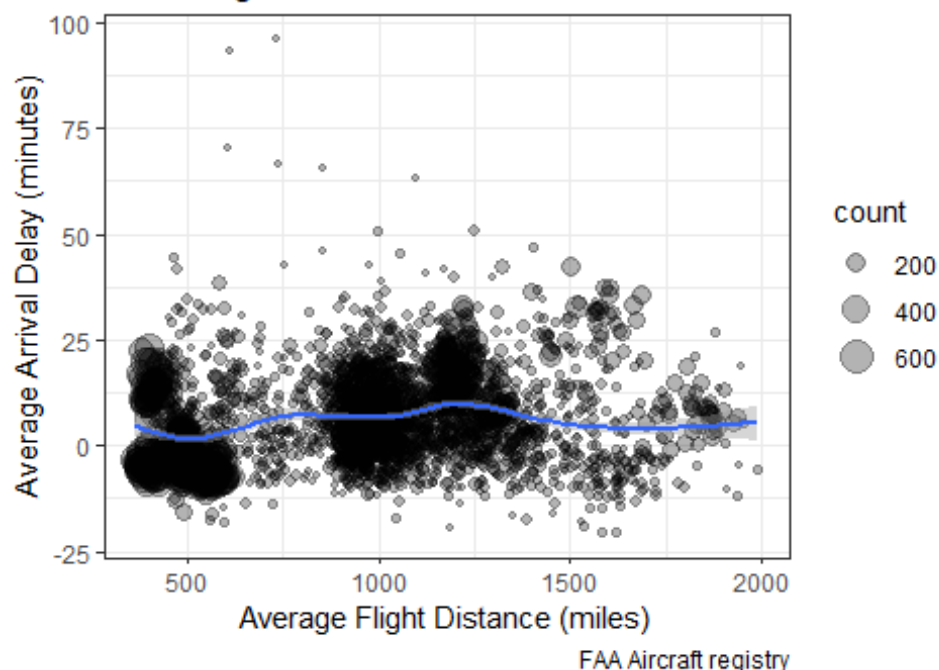
# A tibble: 6 × 4
  tailnum count  dist  delay
  <chr>   <int> <dbl>  <dbl>
1 190NV     29  597.  -9.24
2 191NV      6  583.  -6.67
3 193NV     18  626. -13.8
4 195NV      2  587.  -0.5
5 196NV      3  605. -11.3
6 202NV     17  597. -14.1

delay <- filter(by_tailnum, count > 20, dist < 2000)
# only include counts > 20 and distance < 2000 mi
```

Average delay is only slightly related to average distance flown by a plane.

```
ggplot(delay, aes(dist, delay)) +
  geom_point(aes(size = count), alpha = .3) +
  geom_smooth() +
  scale_size_area() +
  theme_bw() +
  labs(x = "Average Flight Distance (miles)",
       y = "Average Arrival Delay (minutes)",
       caption = "FAA Aircraft registry",
       title = "Flight Distance and Average Arrival Delays \n from Flights from NY")
`geom_smooth()` using method = 'gam' and formula = 'y ~ s(x, bs = "cs")'
```

Flight Distance and Average Arrival Delays from Flights from NY



Heat Tree Stream Alluvial - Part 2

R Saidi

Streamgraphs

```
library(treemap)
library(tidyverse)
library(RColorBrewer)
library(webshot2)
```

Streamgraphs (unfortunately do not render to rpubs)

This type of visualisation is a variation of a stacked area graph, but instead of plotting values against a fixed, straight axis, a streamgraph has values displaced around a varying central baseline. Streamgraphs display the changes in data over time of different categories through the use of flowing, organic shapes that somewhat resemble a river-like stream. This makes streamgraphs aesthetically pleasing and more engaging to look at.

The size of each individual stream shape is proportional to the values in each category. The axis that a streamgraph flows parallel to is used for the timescale. Color can be used to either distinguish each category or to visualize each category's additional quantitative values through varying the color shade.

What are streamgraphs good for?

Streamgraphs are ideal for displaying high-volume datasets, in order to discover trends and patterns over time across a wide range of categories. For example, seasonal peaks and troughs in the stream shape can suggest a periodic pattern. A streamgraph could also be used to visualize the volatility for a large group of assets over a certain period of time.

The downside to a streamgraph is that they suffer from legibility issues, as they are often very cluttered. The categories with smaller values are often drowned out to make way for categories with much larger values, making it impossible to see all the data. Also, it's impossible to read the exact values visualized, as there is no axis to use as a reference.

Streamgraph code

The code for making streamgraphs has changed with new updates to R. You have to download and install Rtools40 from the link, <https://cran.rstudio.com/bin/windows/Rtools/>. and then used the code provided below.

Load devtools and libraries to create the following streamgraphs

install the package “devtools” also run the line: `devtools::install_github("hrbrmstr/streamgraph")` , then comment it out.

```
#install "devtools" (as a package)
devtools::install_github("hrbrmstr/streamgraph")
library(streamgraph) # install "streamgraph" as a package
library(babynames) # install "babynames"
data(babynames)
```

Now look at the babynames dataset

```
ncol(babynames)

[1] 5

head(babynames)

# A tibble: 6 × 5
  year sex  name      n  prop
<dbl> <chr> <chr>   <int> <dbl>
1  1880 F    Mary    7065 0.0724
2  1880 F    Anna    2604 0.0267
3  1880 F    Emma    2003 0.0205
4  1880 F  Elizabeth 1939 0.0199
5  1880 F   Minnie   1746 0.0179
6  1880 F  Margaret 1578 0.0162

summary(babynames$year)

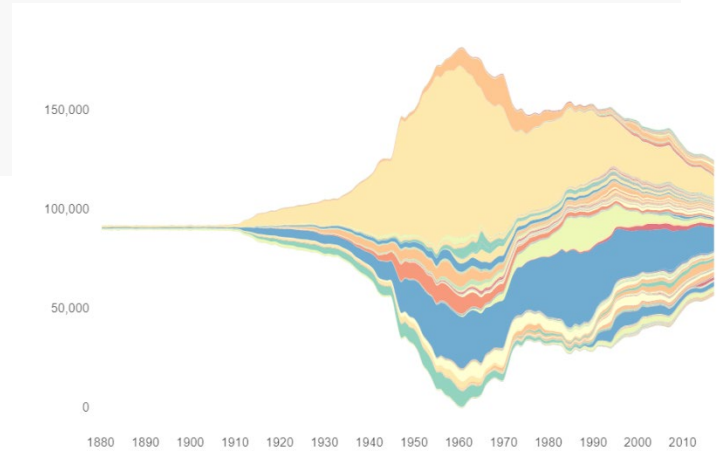
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
  1880   1951   1985   1975   2003   2017
```

Babynames streamgraph

Mouse over the colors and years to look at the pattern of various names

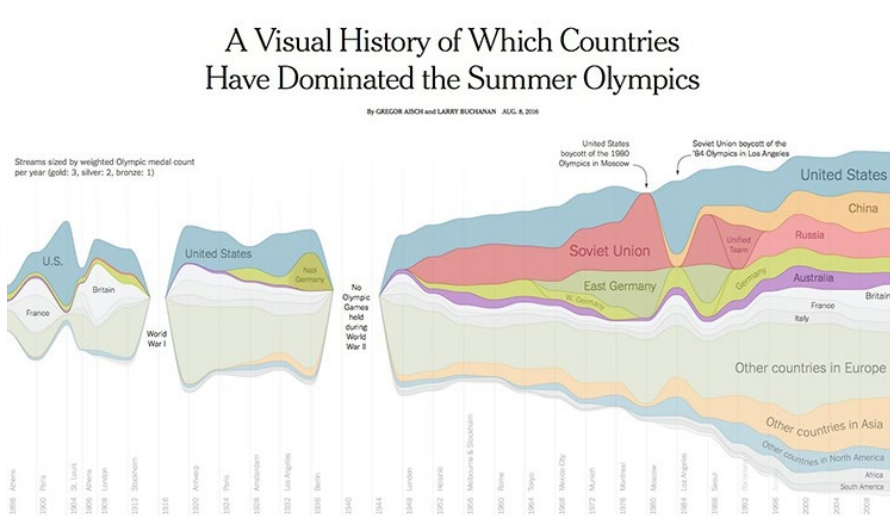
```
babynames |>
  filter(grepl("^Jo", name)) |>
  group_by(year, name) |>
  tally(wt=n) |>
  streamgraph("name", "n", "year")
```

```
babynames |>
  filter(grepl("^Da", name)) |>
  group_by(year, name) |>
  tally(wt=n) |>
  streamgraph("name", "n", "year")
```



Learn from the experts (a streamgraph)

Over the coming weeks and beyond, make a habit of looking for innovative graphics, especially those employing unusual chart forms that communicate the story from data in an effective way. Work out how they use visual cues to encode data. Here are a couple of examples from *The New York Times* to get you started. Follow the links from the source credits to explore the interactive versions:



(Source: *The New York Times*)

Your Assignment This Week 5

1. **(Ungraded)** Copy the Markdown code from these notes to explore how to create treemaps, heatmaps, streamgraphs, and alluvials. You can publish your RMD or Quarto file in Rpubs.
2. **(Worth up to 15 points)** NYC Flights Homework

This week, you will create your first visualization on your own using the pre-built dataset, `nycflights23`. Load the libraries and view the “flights” dataset

```
library(tidyverse)
library(nycflights23)
```

Now create one data visualization with this dataset. Your assignment is to create one plot to visualize one aspect of this dataset. The plot may be any type we have covered so far in this class (bargraphs, scatterplots, boxplots, histograms, treemaps, heatmaps, streamgraphs, or alluvials)

Requirements for the plot:

- a. Include at least one dplyr command (filter, sort, summarize, group_by, select, mutate,)
- b. Include labels for the x- and y-axes
- c. Include a title and caption for the data source
- d. Your plot must incorporate at least 2 colors
- e. Include a legend that indicates what the colors represent
- f. Write a brief paragraph that describes the visualization you have created and at least one aspect of the plot that you would like to highlight.

Start early so that if you do have trouble, you can email me with questions.

Submit this assignment in the assignment dropbox by **11:59 pm** on _____. *You will present in class next week.*